

# Kubernetes有状态集群服务 部署与管理

时速云：张寿红



促进软件开发领域知识与创新的传播



关注InfoQ官方微信  
及时获取ArchSummit  
大会演讲视频信息



全球软件开发大会 [北京站]

2017年4月16-18日 北京·国家会议中心

咨询热线: 010-64738142



全球架构师峰会 2016 [深圳站]

2017年7月7-8日 深圳·华侨城洲际酒店

咨询热线: 010-89880682

# Agenda

- **Background**
  - What is Kubernetes?
  - Run stateful services on Kubernetes
- **Kubernetes Stateful Service Features**
  - Init Container
  - Pet Set
- **Kubernetes Storage**
  - Volume
  - Persistent Volume
  - Dynamic Storage Provision
- **Run MySQL Cluster on Kubernetes**
  - Galera MySQL Introduction
  - Deploy MySQL Cluster with PetSet
  - Cluster Operation on Kubernetes

# What is Kubernetes

Greek for “*Helmsman*”; also the root of the words “*governor*” and “*cybernetic*”

Runs and manages containers

Handles failure

Inspired and informed by Google’s experiences and internal systems

Supports multiple cloud and bare-metal environments

Supports multiple container runtimes

**100% Open source**, part of CNCF, written in Go

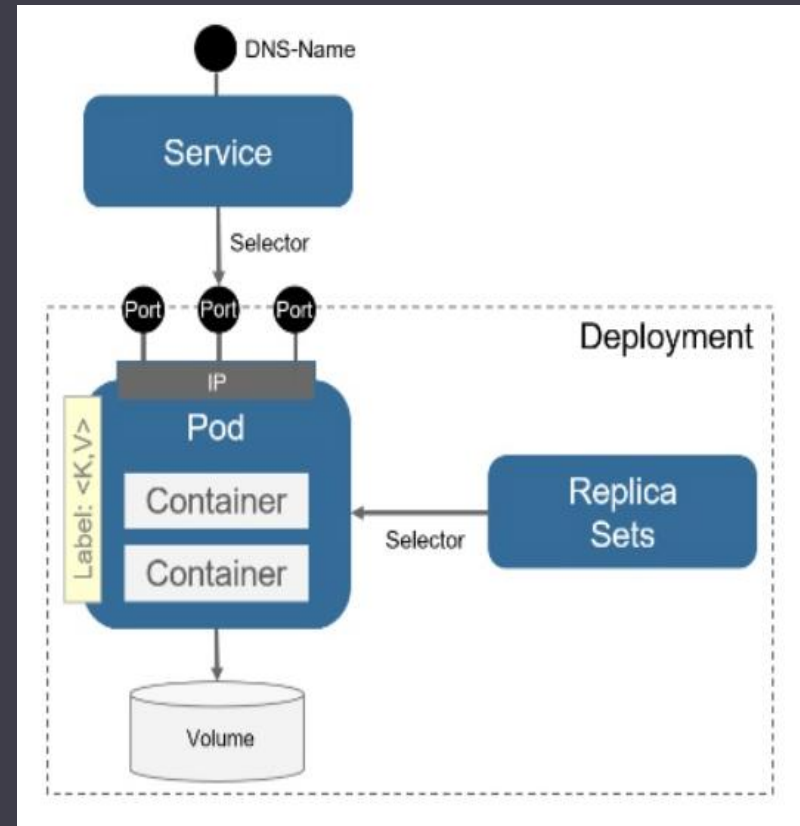


docker

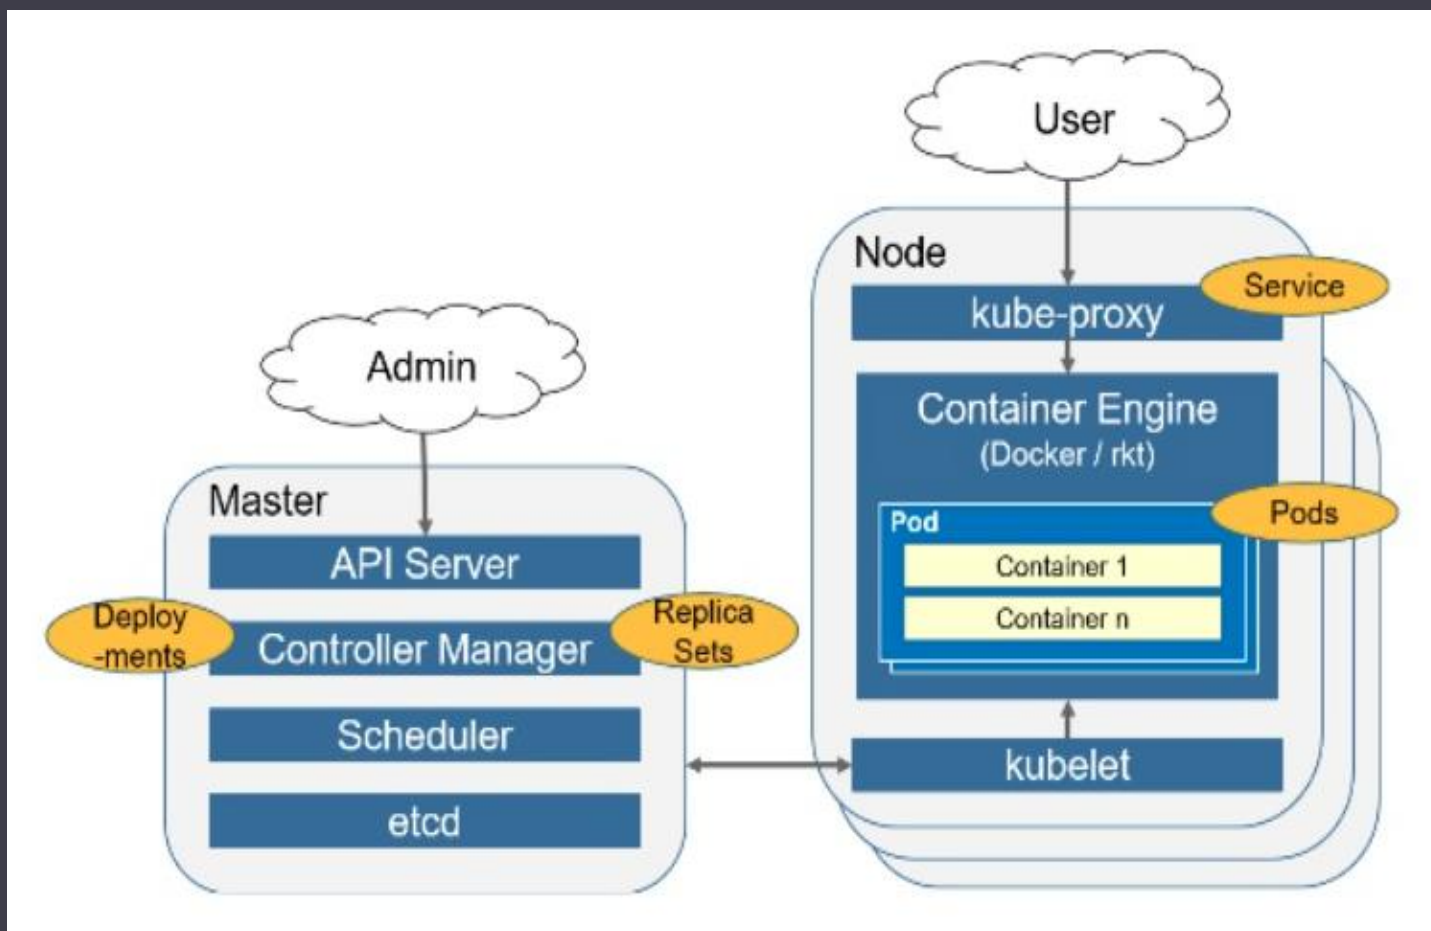


# What is Kubernetes

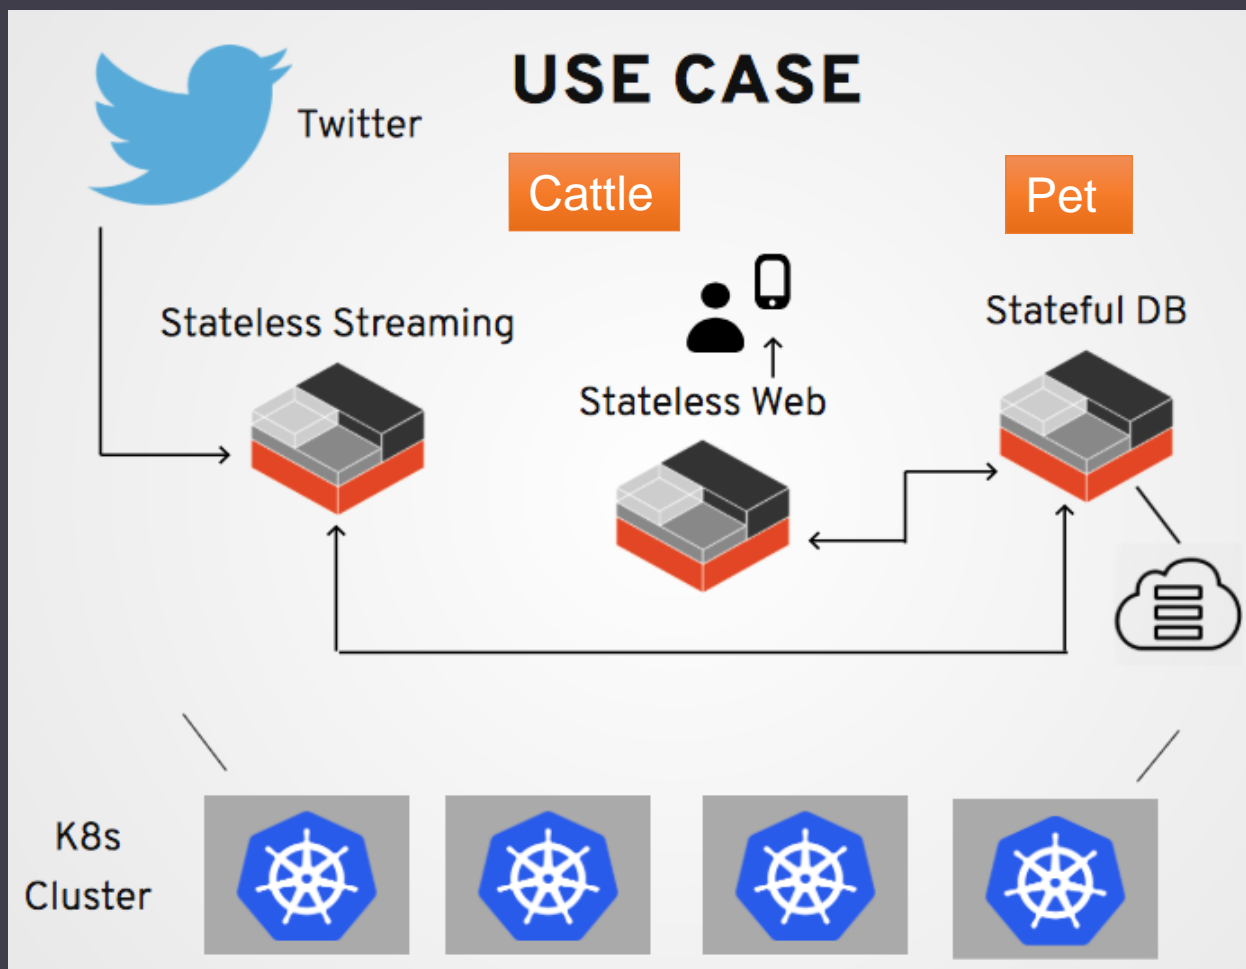
- **Services** are an abstraction for a logical set of Pods.
- **Pods** are the smallest deployable units of computing.
- **Deployments** provide declarative updates for Pods and RCs.
- **Replica Sets** ensure specified number of Pods are running.
- **Labels** are key/value pairs attached to objects used for identification.



# What is Kubernetes



# What is Kubernetes



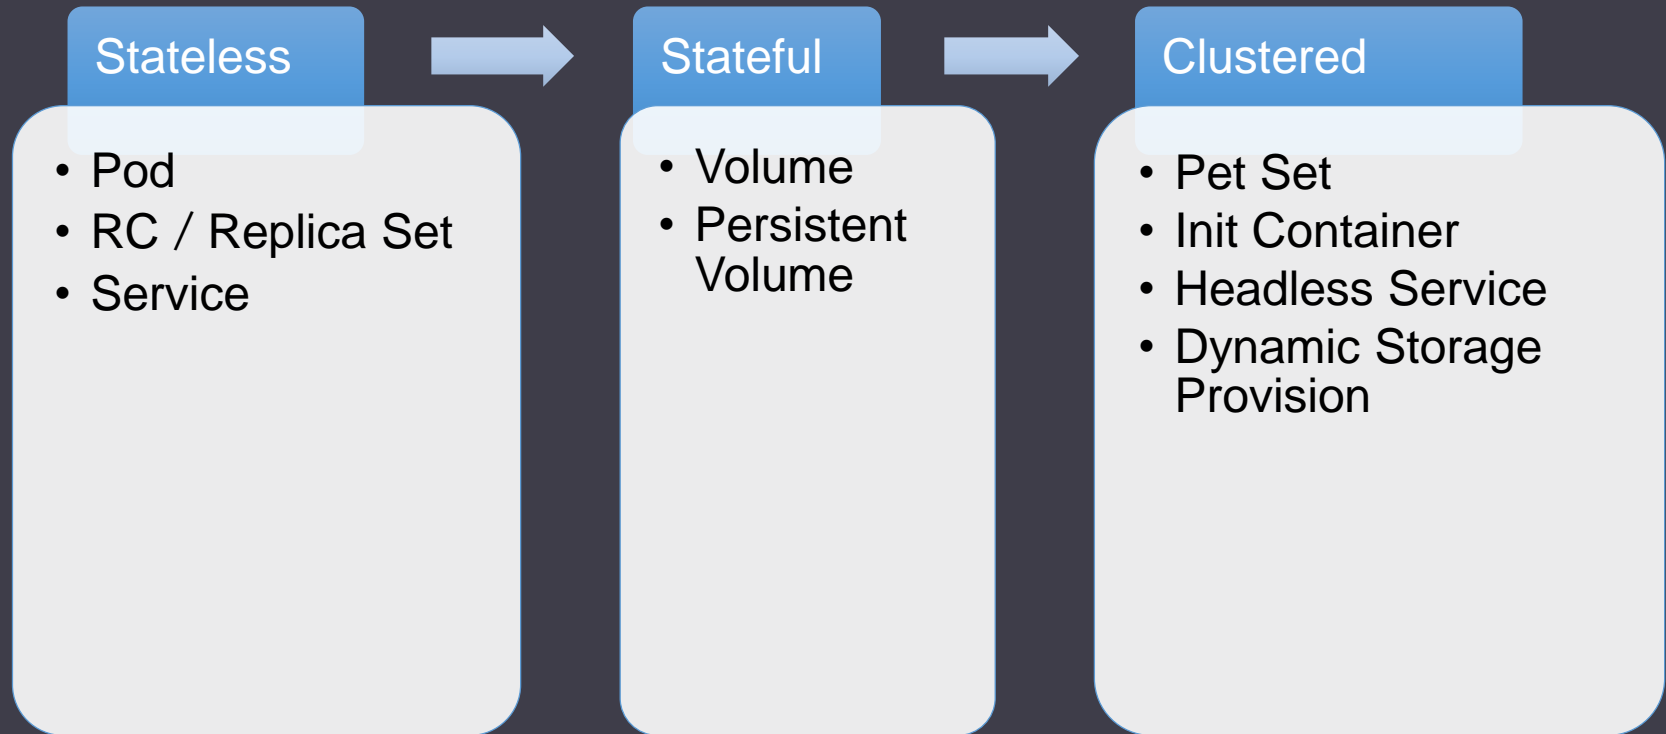


# Run Stateful Service on Kubernetes

- Stateless services are popular in containerized world because they are disposable and easy to be re-created with container images.
- Beyond stateless services like web server, users are increasingly deploying stateful services with containers to benefit from “build once, run anywhere” and to improve bare metal efficiency/utilization.
- These “pets” bring new requirements including longer life cycle, configuration dependencies and stateful failover. Container orchestration must address these needs to successfully deploy and scale apps.



# Run Stateful Service on Kubernetes

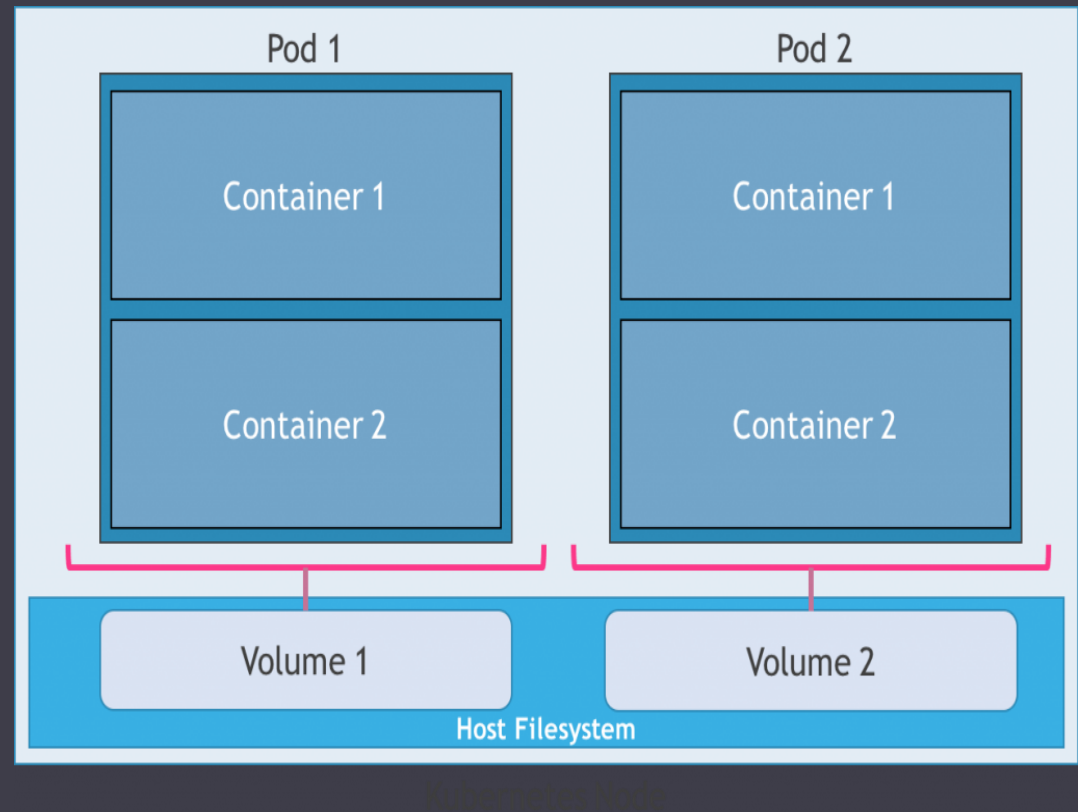


# Agenda

- **Background**
  - What is Kubernetes?
  - Run stateful services on Kubernetes
- **Kubernetes Storage**
  - Volume
  - Persistent Volume
  - Dynamic Storage Provision
- **Kubernetes Stateful Service Features**
  - Init Container
  - Pet Set
- **Run MySQL Cluster on Kubernetes**
  - Galera MySQL Introduction
  - Deploy MySQL Cluster with PetSet
  - Cluster Operation on Kubernetes

# Kubernetes Storage


- Node Based Volume
  - emptyDir
  - hostPath
- Use Cases:
  - Store temp data
  - Data sharing among containers



# Kubernetes Storage


- emptyDir

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-persistent-storage
          mountPath: /data/redis
      volumes:
        - name: redis-persistent-storage
          emptyDir: {}
```



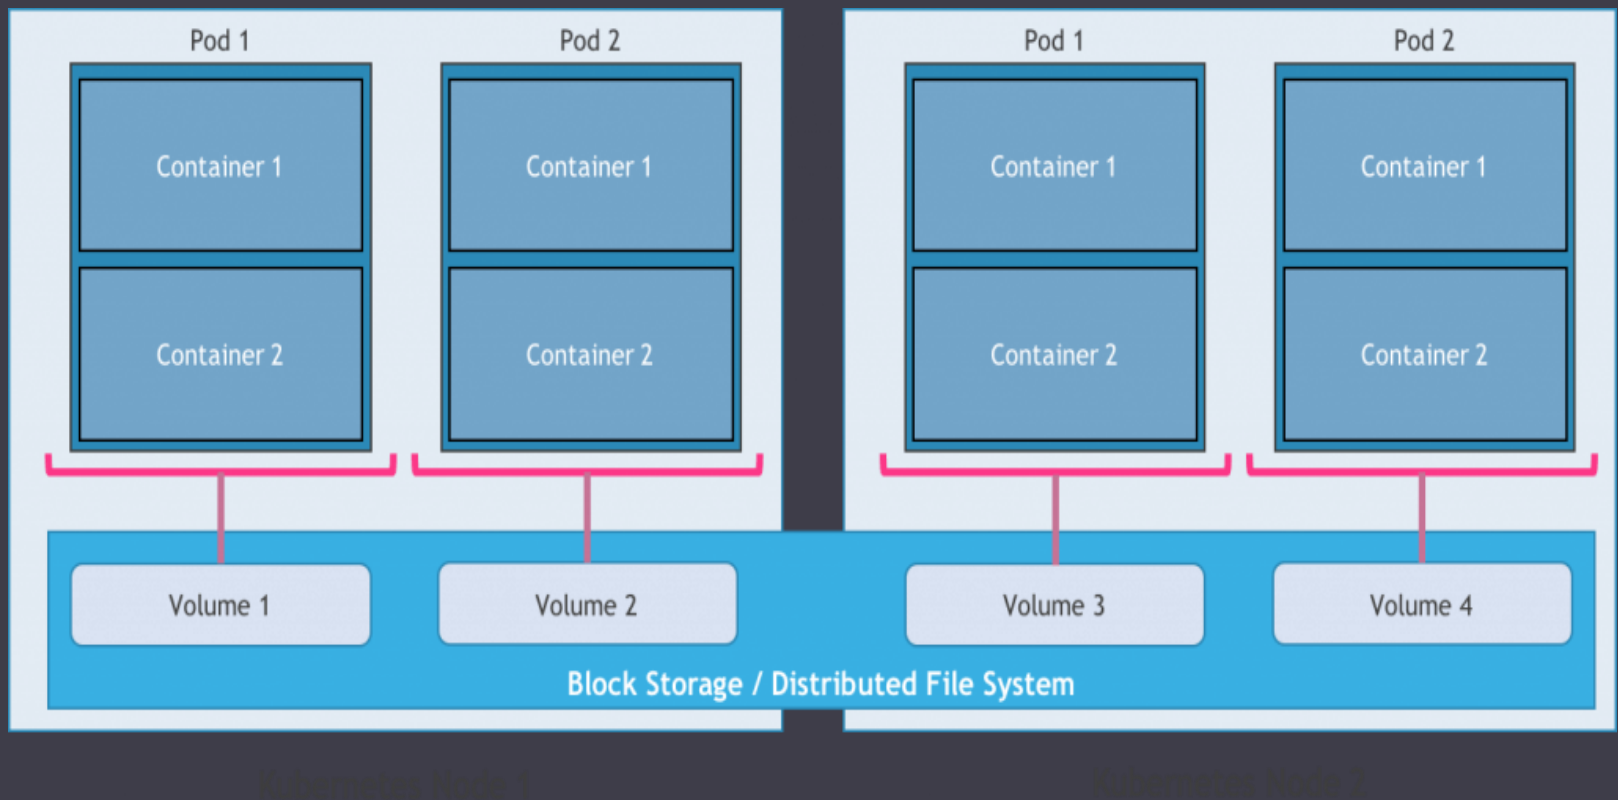
- hostPath

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pd
spec:
  containers:
    - image: gcr.io/google_containers/test-webserver
      name: test-container
      volumeMounts:
        - mountPath: /test-pd
          name: test-volume
      volumes:
        - name: test-volume
          hostPath:
            # directory location on host
            path: /data
```




# Kubernetes Storage

- Cross Node Volume



# Kubernetes Storage



Various volume plugins

Shared Volume

Unshared Volume

- nfs
- hostPath
- glusterfs
- flocker
- cephfs
- secret
- vsphereVirtualDisk
- emptyDir
- hostPath
- gcePersistentDisk
- awsElasticBlockStore
- iscsi
- rbd
- azureFileVolume

# Kubernetes Storage


- Add your volume plugin

1. Register in kubelet entry  
kubelet/app/plugins.go

```
allPlugins = append(allPlugins, customizedPlugin.ProbeVolumePlugins()...)
```

2. Implement it in the package below:

pkg/volume/<your\_plugin>, interface can refer to volume.go



- ProbeVolumePlugins
- Init
- CanSupport
- SetUp
- TearDown
- ...

3. Update API Spec



# Kubernetes Storage

- Persistent Volume & Persistent Volume Claim
  - PV/PVC abstracts details of how storage is provided from how it is consumed.
  - PV/PVC are API resources.
  - PVs are volume plugins like Volumes, but have a lifecycle independent of any individual pod that uses the PV.

# Kubernetes Storage

- PV Access Modes

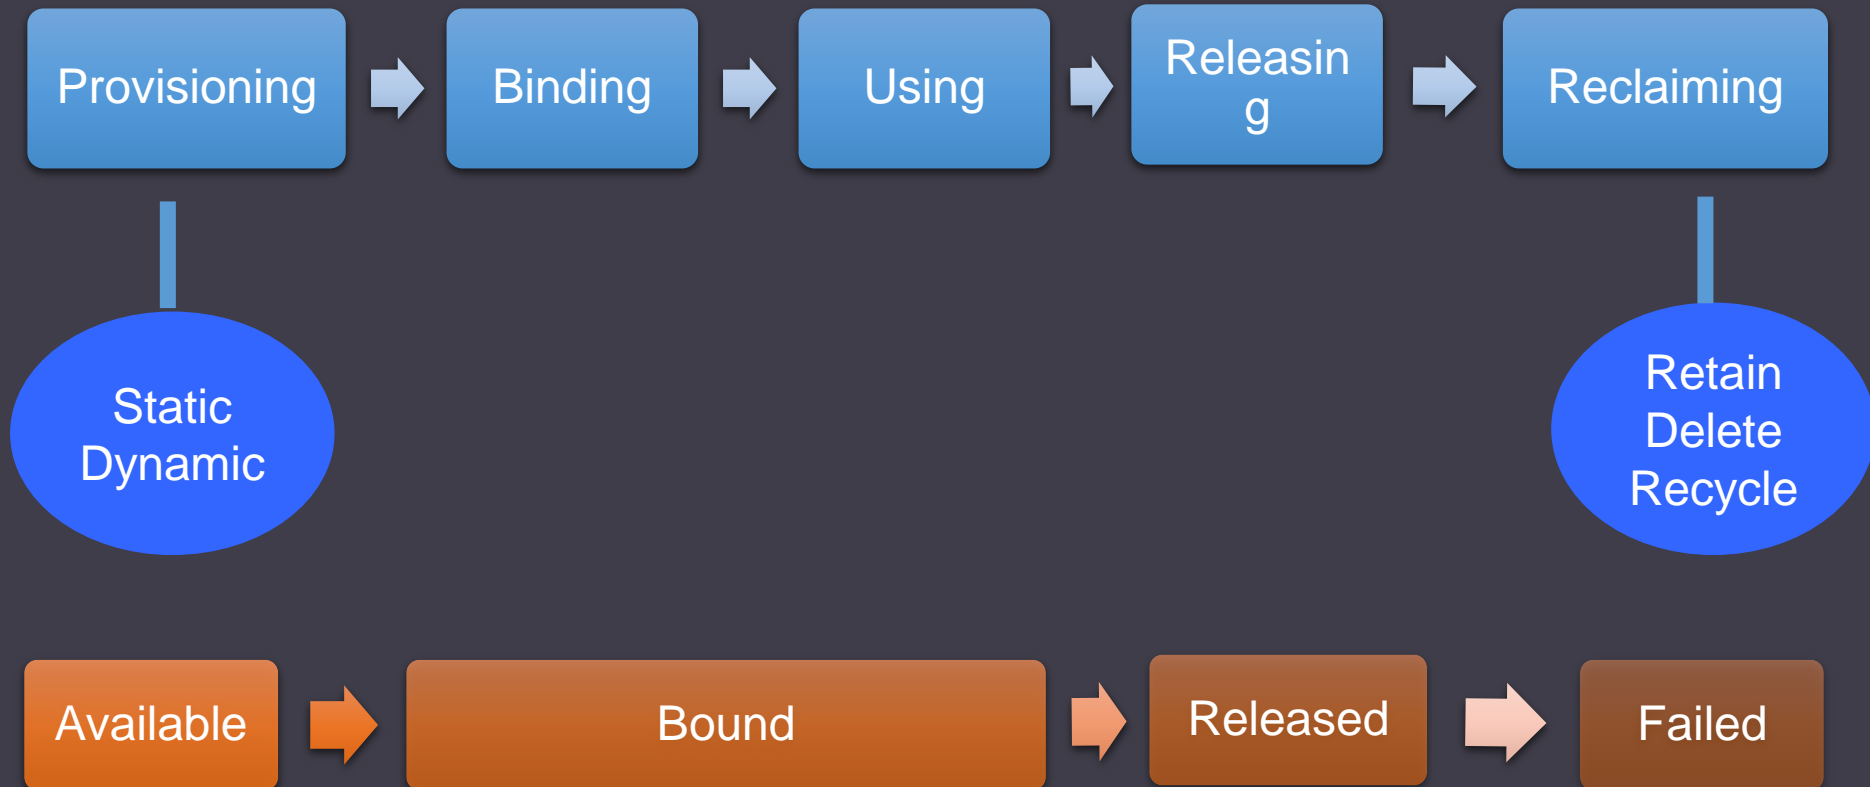
Volume Plugin	ReadWriteOnce	ReadOnlyMany	ReadWriteMany
AWSElasticBlockStore	x	-	-
AzureFile	x	x	x
CephFS	x	x	x
Cinder	x	-	-
FC	x	x	-
FlexVolume	x	x	-
GCEPersistentDisk	x	x	-
Glusterfs	x	x	x
HostPath	x	-	-
iSCSI	x	x	-
NFS	x	x	x
RDB	x	x	-
VsphereVolume	x	-	-

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: myclaim
  annotations:
    volume.beta.kubernetes.io/
spec:
  accessModes:
    - ReadWriteOnce
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
  annotations:
    volume.beta.kubernetes.io/
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteOnce
```

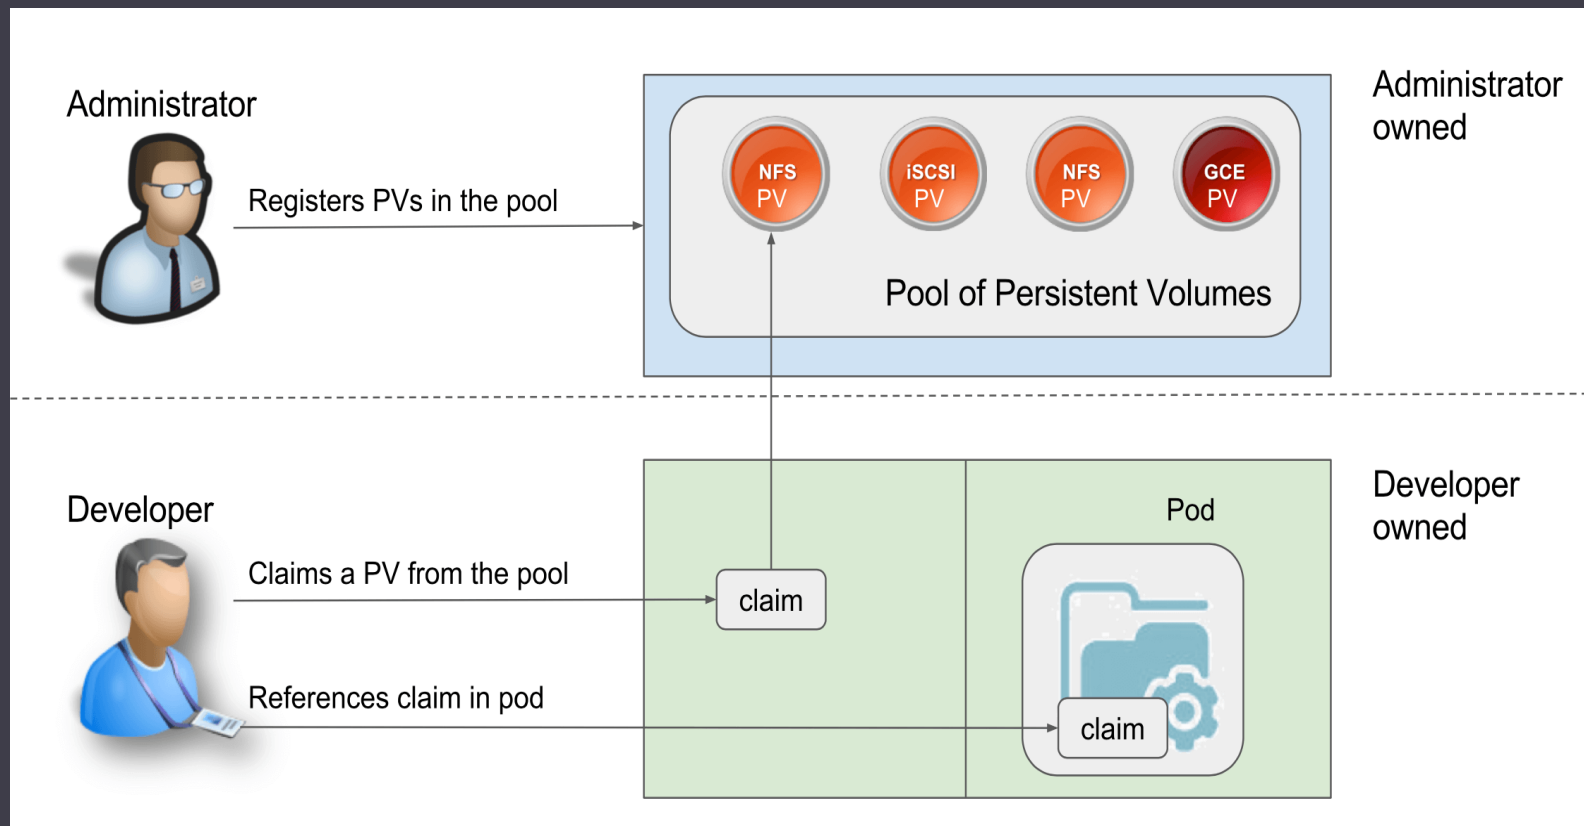
# Kubernetes Storage

- PV Lifecycle



# Kubernetes Storage

- PV Provision & Binding Process



# Kubernetes Storage

- Dynamic Storage Provision

```
kind: StorageClass
```

```
apiVersion:
```

```
storage.k8s.io/v1beta1
```

```
metadata:
```

```
  name: slow
```

```
provisioner: kubernetes.io/gce-  
pd
```

```
parameters:
```

```
  type: pd-standard
```

```
kind: StorageClass
```

```
apiVersion:
```

```
storage.k8s.io/v1beta1
```

```
metadata:
```

```
  name: fast
```

```
provisioner:  
kubernetes.io/gce-pd
```

```
parameters:
```

```
  type: pd-ssd
```

# Kubernetes Storage

- Dynamic Storage Provision

```
"kind": "PersistentVolumeClaim",
"apiVersion": "v1",
"metadata": {
  "name": "claim1",
  "annotations": {

"volume.beta.kubernetes.io/storage-
class": "fast"
  }
},
```

```
"spec": {
  "accessModes": [
    "ReadWriteOnce"
  ],
  "resources": {
    "requests": {
      "storage": "30Gi"
    }
  }
}
```

This claim will result in an SSD-like Persistent Disk being automatically provisioned. When the claim is deleted, the volume will be destroyed.

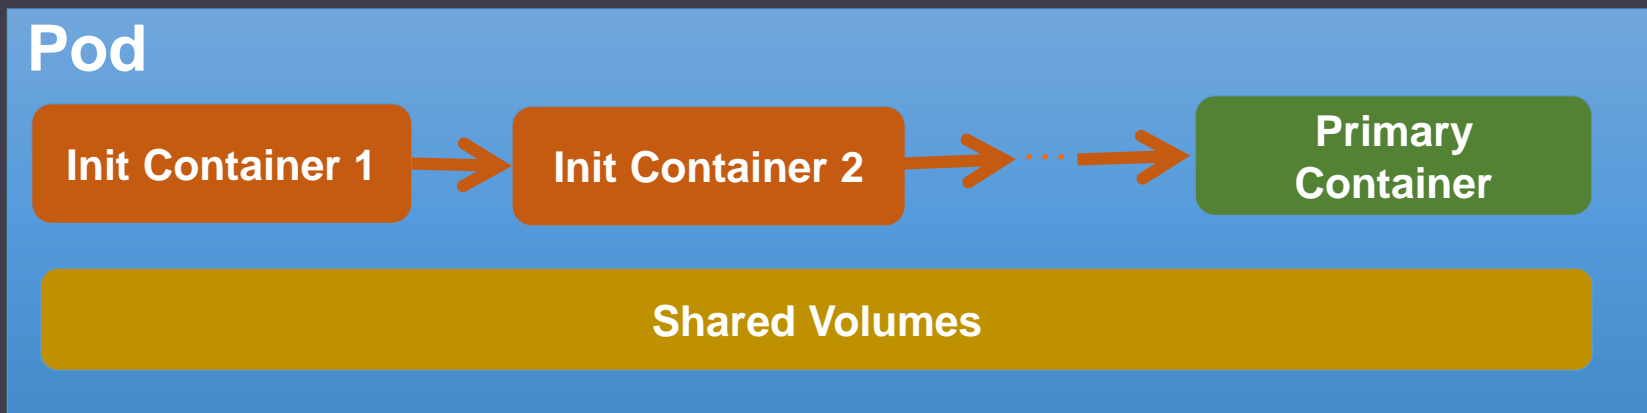
# Agenda

- **Background**
  - What is Kubernetes?
  - Run stateful services on Kubernetes
- **Kubernetes Stateful Service Features**
  - Init Container
  - Pet Set
- **Kubernetes Storage**
  - Volume
  - Persistent Volume
  - Dynamic Storage Provision
- **Run MySQL Cluster on Kubernetes**
  - Galera MySQL Introduction
  - Deploy MySQL Cluster with PetSet
  - Cluster Operation on Kubernetes



# Init Container

- What is Init Container
  - Sequentially executed containers in a pod
  - Initialize shared volumes
  - Always run to completion
  - Used on a pod, replica set, deployment, daemon set, pet set or job.



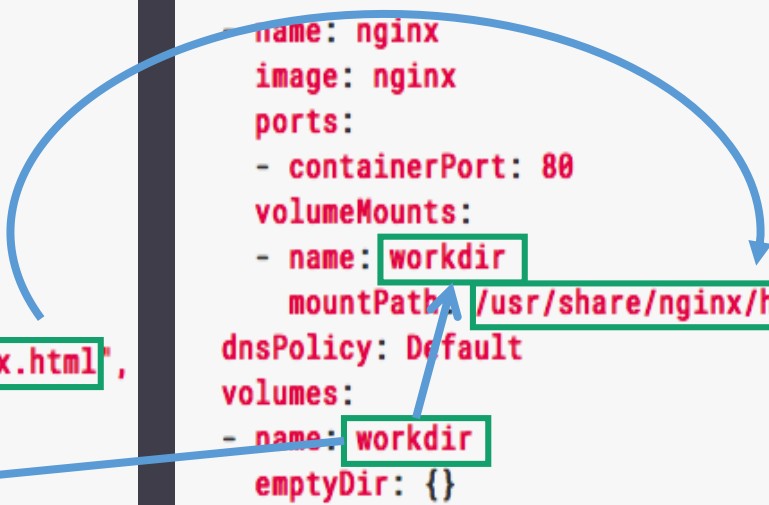
# Init Container

- Why Init Container
  - Waiting for other components to be available
  - Performing configuration
  - Registering the pod into a central database
  - Downloading application dependencies
  - ...

# Init Container

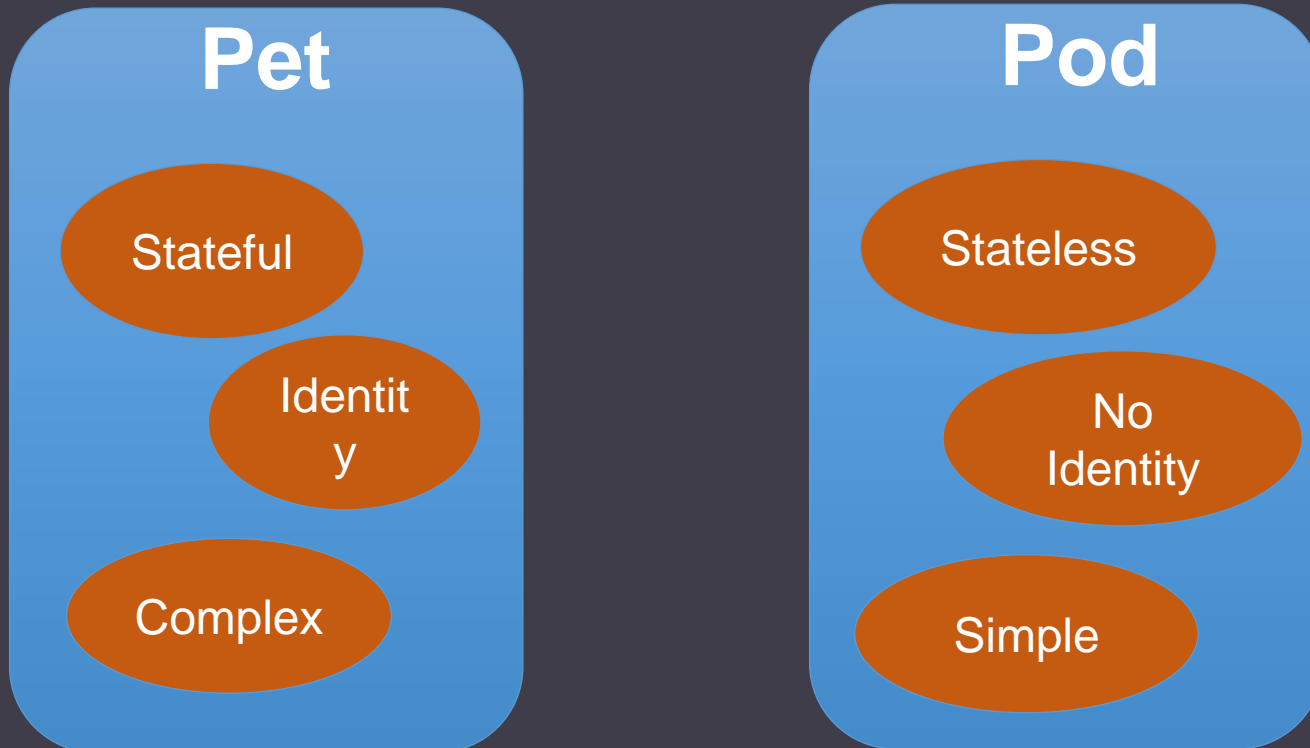
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  annotations:
    pod.alpha.kubernetes.io/init-containers: '[
      {
        "name": "install",
        "image": "busybox",
        "command": ["wget", "-O", "/work-dir/index.html"],
        "volumeMounts": [
          {
            "name": "workdir",
            "mountPath": "/work-dir"
          }
        ]
      }
    ]'
```

```
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
      volumeMounts:
        - name: workdir
          mountPath: /usr/share/nginx/html
      dnsPolicy: Default
  volumes:
    - name: workdir
      emptyDir: {}
```

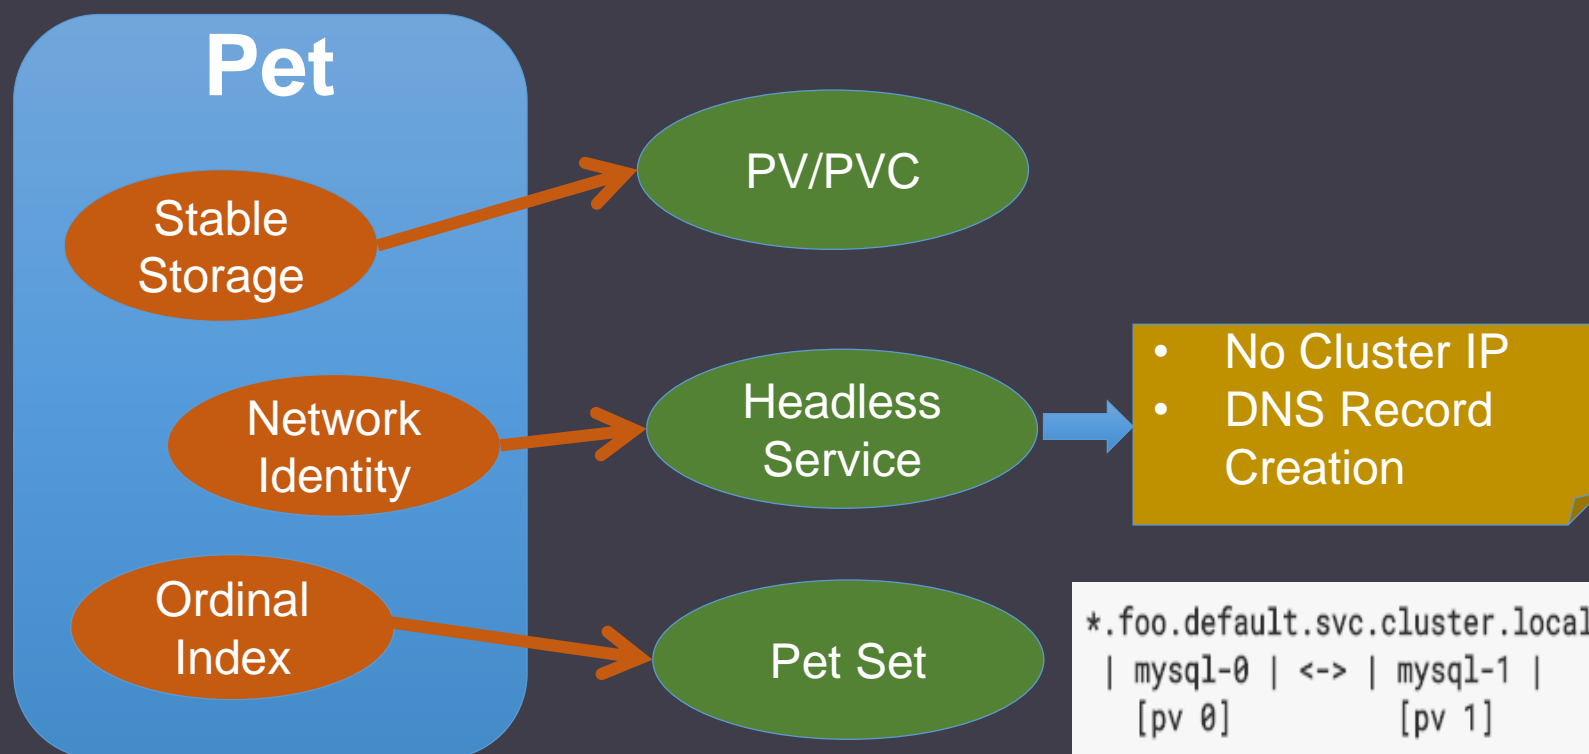


# Pet Set

- Pet vs Pod



# Pet Set



# Pet Set

- PetSet Operations
  - Peer discovery
  - Scaling a PetSet
  - Image upgrades
  - Deleting a PetSet

```
web-0 # nslookup -type=svr nginx.default
Server:      10.0.0.10
Address:     10.0.0.10#53

web-1.nginx.default.svc.cluster.local
web-0.nginx.default.svc.cluster.local
```

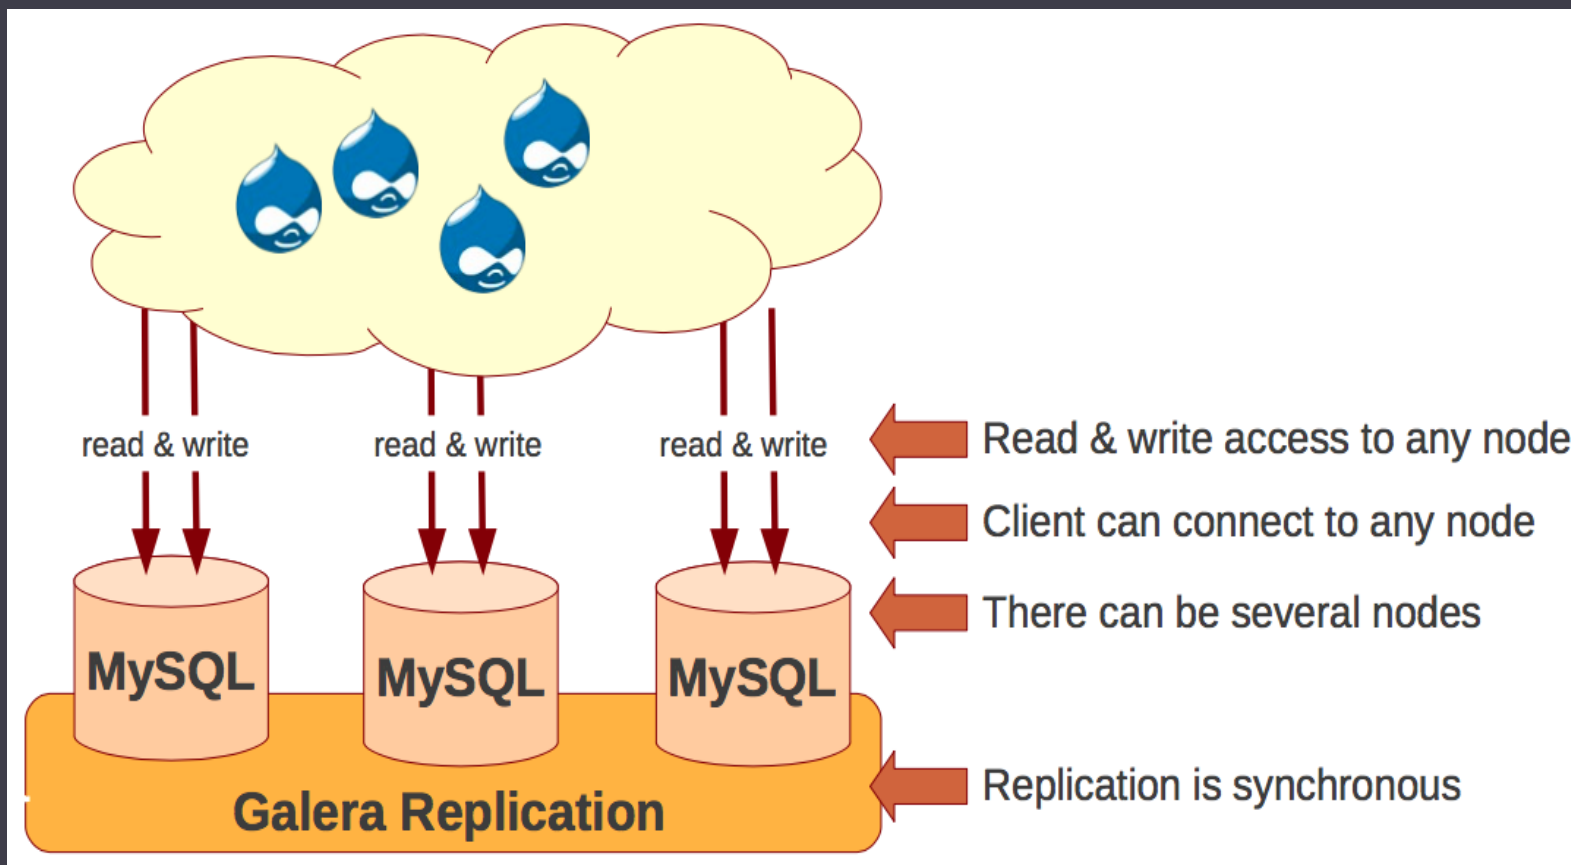
```
$ kubectl scale petset web --replicas=5
petset "web" scaled
```

# Agenda

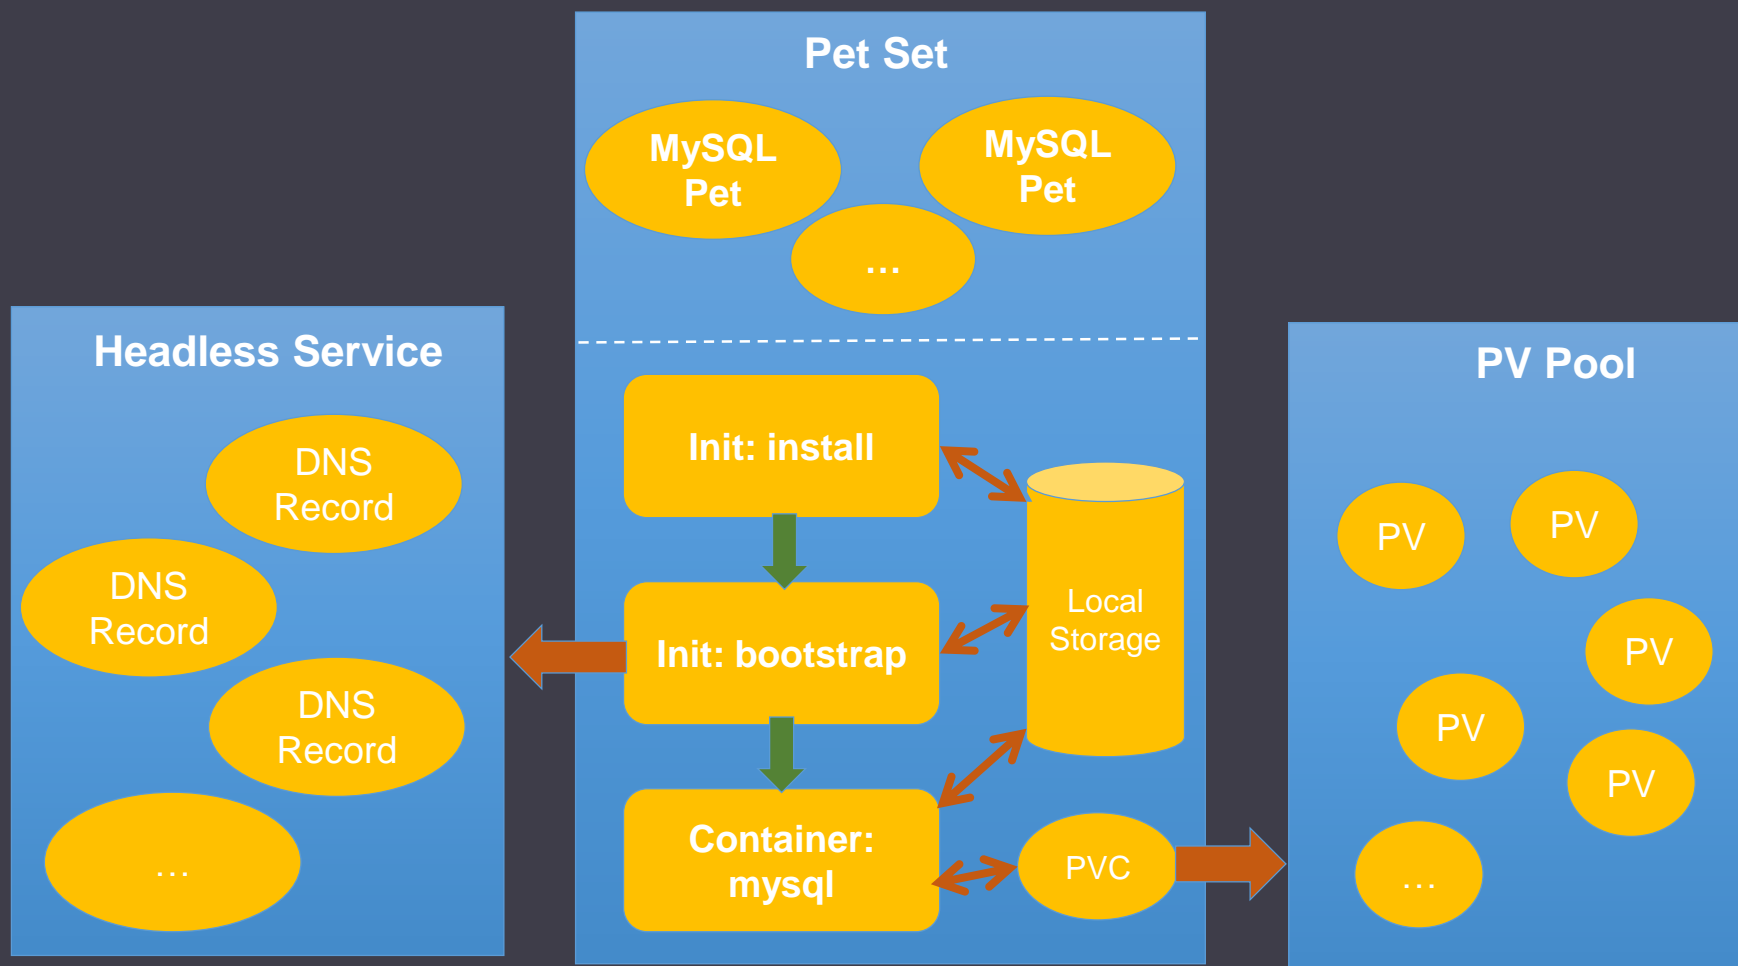
- **Background**
  - What is Kubernetes?
  - Run stateful services on Kubernetes
- **Kubernetes Stateful Service Features**
  - Init Container
  - Pet Set
- **Kubernetes Storage**
  - Volume
  - Persistent Volume
  - Dynamic Storage Provision
- **Run MySQL Cluster on Kubernetes**
  - Galera MySQL Introduction
  - Deploy MySQL Cluster with PetSet
  - Cluster Operation on Kubernetes



# Galera MySQL Introduction



# Deploy MySQL Cluster with PetSet



# Deploy MySQL Cluster with PetSet

## Headless Service & PetSet

```
apiVersion: apps/v1alpha1
kind: PetSet
metadata:
  name: mysql
spec:
  serviceName: "galera"
  replicas: 3
  template:
    metadata:
      labels:
        app: mysql
```

```
# A headless service to create DNS records
apiVersion: v1
kind: Service
metadata:
  annotations:
    service.alpha.kubernetes.io/tolerate-unready-endpoints: "true"
  name: galera
  labels:
    app: mysql
spec:
  ports:
    - port: 3306
      name: mysql
  # *.galeaer.default.svc.cluster.local
  clusterIP: None
  selector:
    app: mysql
```

# Deploy MySQL Cluster with PetSet

## Init Containers

```
annotations:
  pod.alpha.kubernetes.io/initialized: "true"
  pod.alpha.kubernetes.io/init-containers: '[
    {
      "name": "install",
      "image": "gcr.io/google_containers/galera-install:0.1",
      "imagePullPolicy": "Always",
      "args": ["--work-dir=/work-dir"],
      "volumeMounts": [
        {
          "name": "workdir",
          "mountPath": "/work-dir"
        },
        {
          "name": "config",
          "mountPath": "/etc/mysql"
        }
      ]
    }
  ],
```

```
{
  "name": "bootstrap",
  "image": "debian:jessie",
  "command": ["/work-dir/peer-finder"],
  "args": ["-on-start=/work-dir/on-start.sh", "-service=galera"],
  "env": [
    {
      "name": "POD_NAMESPACE",
      "valueFrom": {
        "fieldRef": {
          "apiVersion": "v1",
          "fieldPath": "metadata.namespace"
        }
      }
    }
  ],
  "volumeMounts": [
    {
      "name": "workdir",
      "mountPath": "/work-dir"
    },
    {
      "name": "config",
      "mountPath": "/etc/mysql"
    }
  ]
}
```

# Deploy MySQL Cluster with PetSet

## Galera MySQL Container

```
spec:
  terminationGracePeriodSeconds: 0
  containers:
  - name: mysql
    image: gcr.io/google_containers/mysql-galera:e2e
    ports:
    - containerPort: 3306
      name: mysql
    - containerPort: 4444
      name: sst
    - containerPort: 4567
      name: replication
    - containerPort: 4568
      name: ist
```

```
args:
  - --defaults-file=/etc/mysql/my-galera.cnf
  - --user=root
readinessProbe:
  # TODO: If docker exec is buggy just use g
  exec:
    command:
    - sh
    - -c
    - "mysql -u root -e 'show databases;'"
  initialDelaySeconds: 15
  timeoutSeconds: 5
  successThreshold: 2
volumeMounts:
  - name: datadir
    mountPath: /var/lib/
  - name: config
    mountPath: /etc/mysql
```

# Deploy MySQL Cluster with PetSet

## Volumes & PVC

```
volumes:
- name: config
  emptyDir: {}
- name: workdir
  emptyDir: {}
volumeClaimTemplates:
- metadata:
  name: datadir
  annotations:
    volume.alpha.kubernetes.io/storage-class: anything
  spec:
    accessModes: [ "ReadWriteOnce" ]
    resources:
      requests:
        storage: 1Gi
```

## Create Pet Set

```
$ kubectl create -f petset.yaml
```

# Cluster Operation

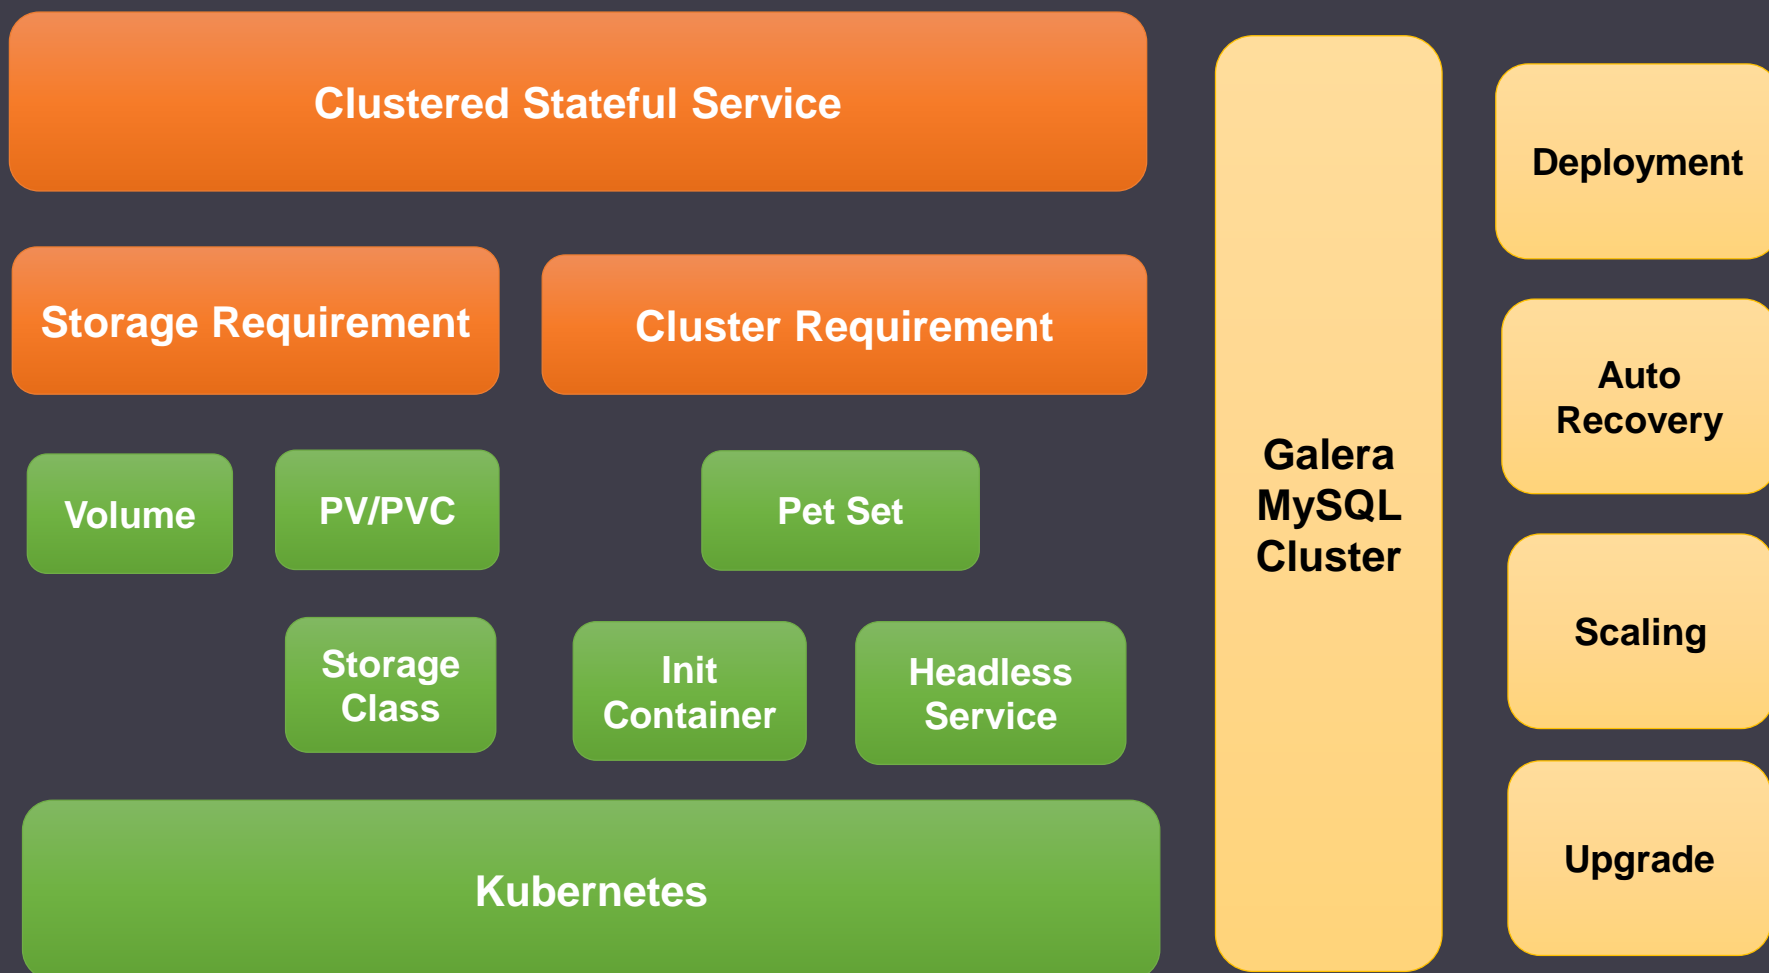
- Auto Recovery
  - The failed pets can be automatically recreated
  - The new created pet will use the data of the died one
- Scaling the Cluster
  - `kubectl scale petset mysql --replicas=5`
- Image Upgrades
  - Update the image field of any container in the podTemplate
  - Delete Pets one by one, the PetSet controller will recreate it with the new image



# Agenda

- **Background**
  - What is Kubernetes?
  - Run stateful services on Kubernetes
- **Kubernetes Stateful Service Features**
  - Init Container
  - Pet Set
- **Kubernetes Storage**
  - Volume
  - Persistent Volume
  - Dynamic Storage Provision
- **Run MySQL Cluster on Kubernetes**
  - Galera MySQL Introduction
  - Deploy MySQL Cluster with PetSet
  - Cluster Operation on Kubernetes

# Summary



# Thanks !

